

Overview of How CNNs, RNNs, and GANs Apply to Real Tasks

Rushi Patel (12202120601038) | Deep Learning and Applications (202047804)

Sem 7 (Artificial Intelligence and Data Science)



What is Deep Learning?

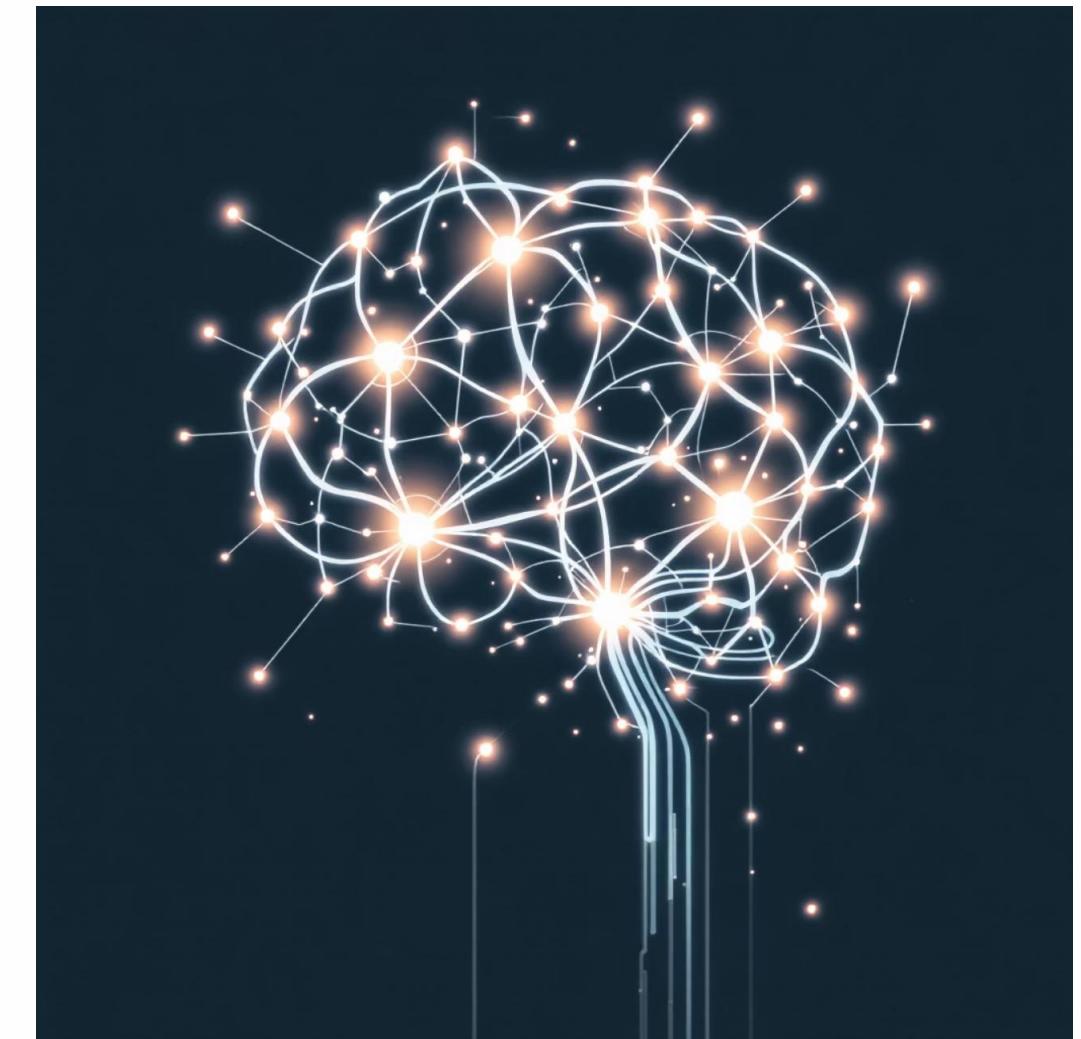
Deep Learning is a specialized subfield of Machine Learning that uses artificial neural networks with multiple layers (hence "deep") to analyze raw data and extract features automatically. This architecture allows it to model complex patterns in data with minimal human intervention.

Traditional ML

Requires manual feature extraction and engineering.

Deep Learning

Learns features directly from raw data, enhancing scalability.



From Simple ANNs to Advanced Architectures

Artificial Neural Networks (ANNs) mimic biological neurons to process information. However, simple, fully connected ANNs struggle with highly structured data like high-resolution images or long text sequences.

ANNs (Basic)

Good for simple classification tasks, but lack spatial or temporal awareness.

Motivation for Advancement

Need to process complex data types efficiently (images, sequences, generative tasks).

Advanced Models

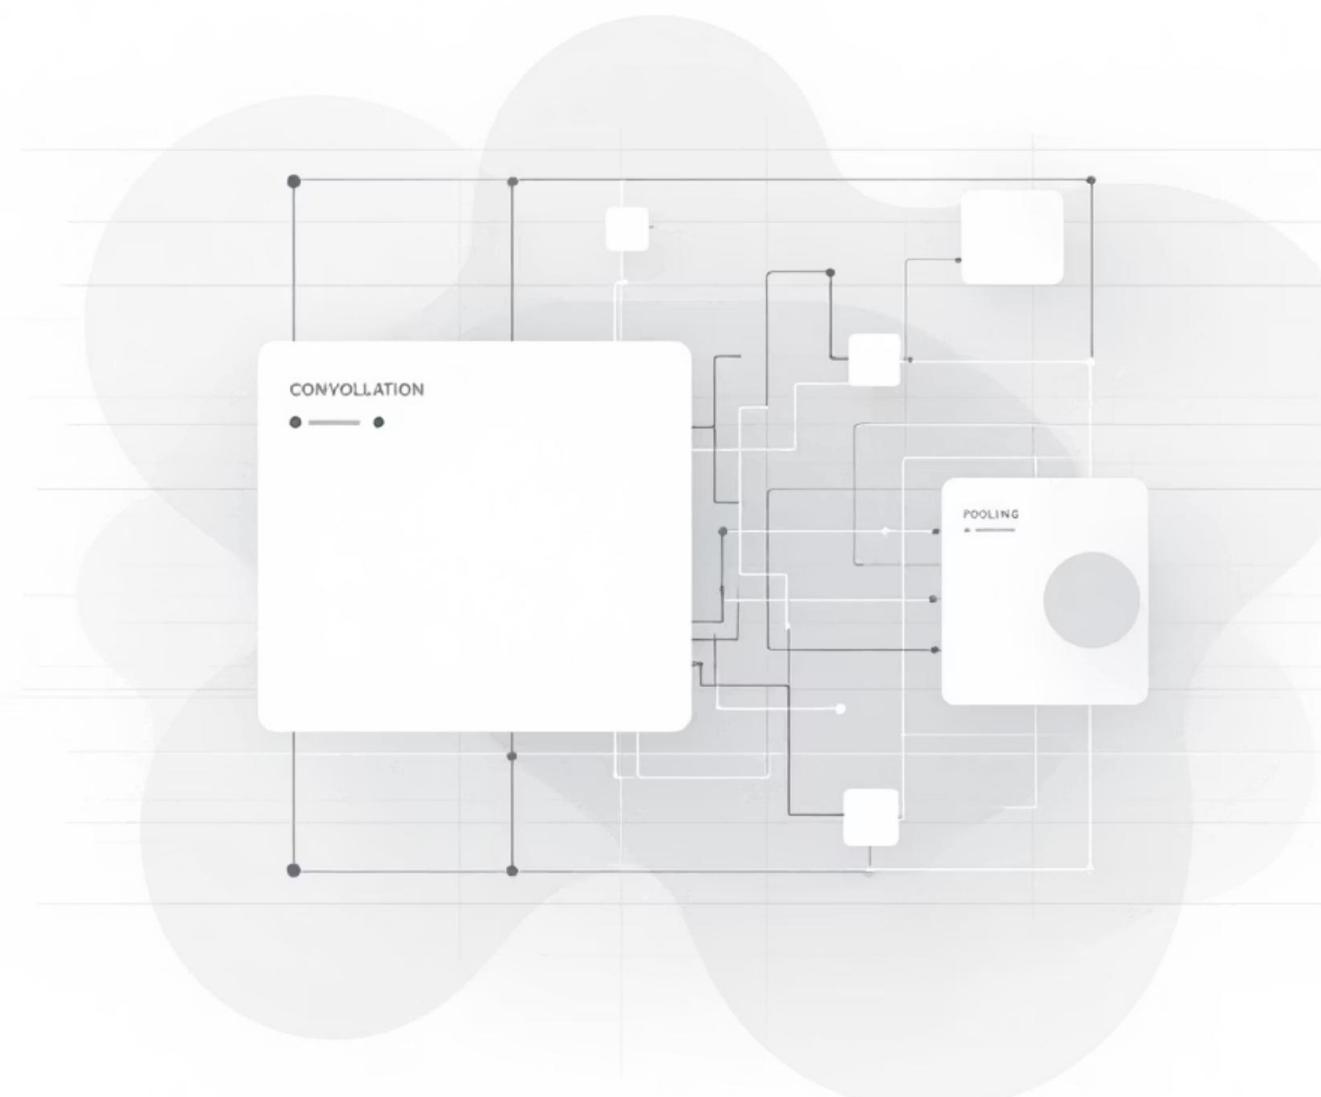
Led to specialized networks: CNNs (Spatial), RNNs (Temporal), GANs (Generative).

Convolutional Neural Networks (CNNs)

CNNs are the workhorses of computer vision, specifically designed to process data with a known spatial relationship, like images and video frames. They achieve translation invariance by using shared weights across local receptive fields.

Core Components:

- **Convolutional Layer:** Learns features (edges, textures) through kernels.
- **Pooling Layer:** Reduces dimensionality and computation, making the model more robust.
- **Fully Connected Layer:** Performs final classification based on extracted features.



Key Applications of CNNs in the Real World



Object Detection

Systems like YOLO and Faster R-CNN use CNNs to locate and classify multiple objects within an image, crucial for robotics and surveillance.



Medical Diagnostics

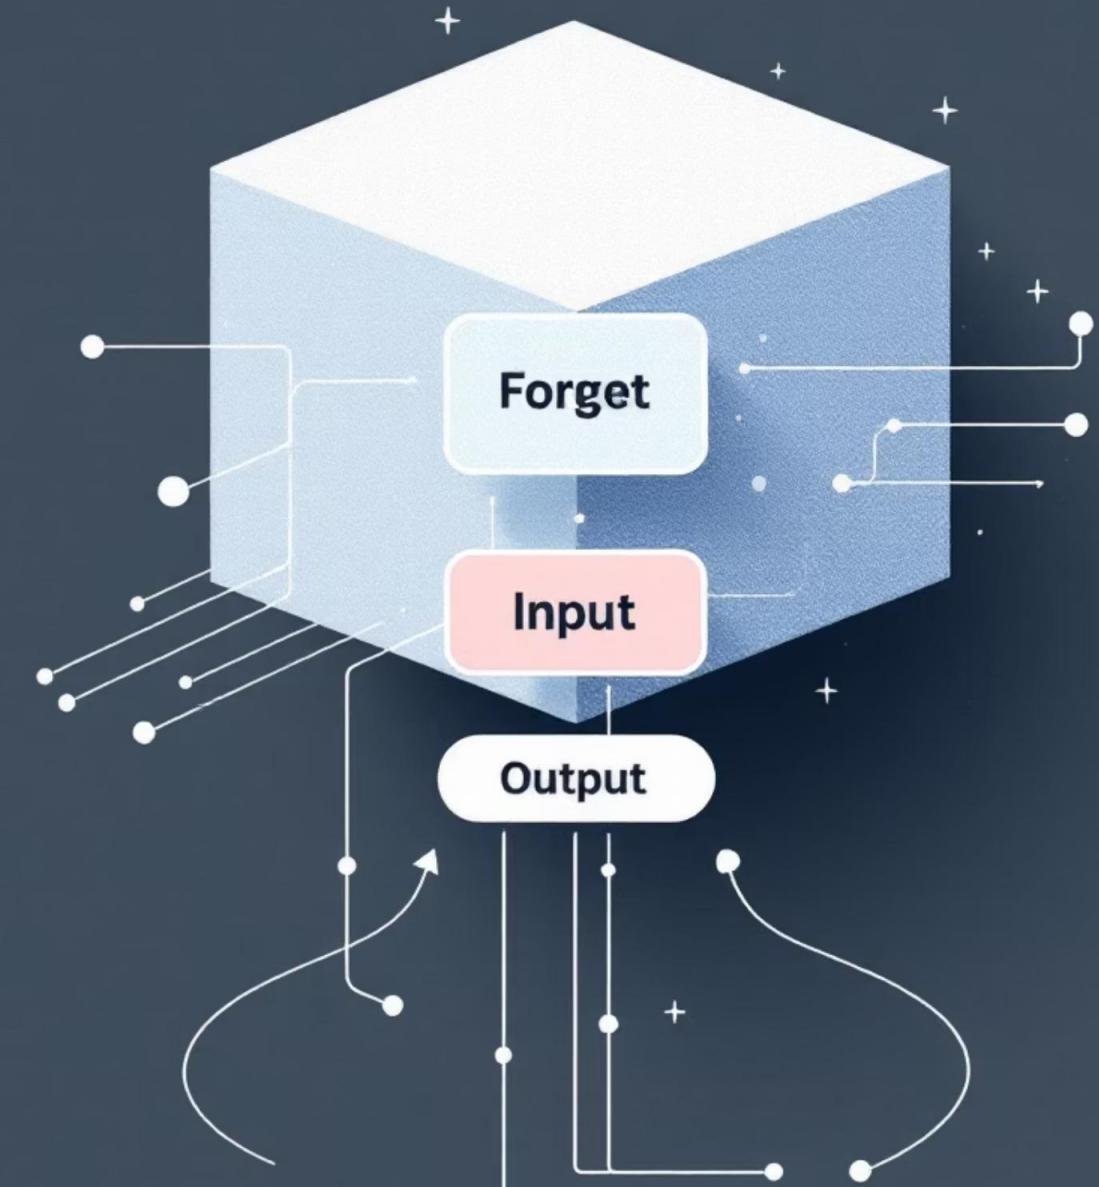
CNNs analyze complex medical scans (X-rays, MRIs) to assist doctors in rapidly detecting diseases like cancer or pneumonia with high accuracy.



Autonomous Vehicles

They process real-time video feed to recognize traffic signs, pedestrians, lanes, and obstacles, ensuring safe navigation.

Recurrent Neural Networks (RNNs)



LSTM network can と、劣い蘊藏性をかない、
LSTM は遡りのいとおしのく、

- 1 Sequential Input
RNNs process sequences of data (text, audio, time series) by looping information from one step to the next.
- 2 Internal Memory
The hidden state acts as memory, allowing the network to use past context to inform current predictions.
- 3 Problem Solved
Long-term dependencies (Vanishing/Exploding Gradient) are mitigated by advanced variants like LSTMs and GRUs.

RNNs in Action: Understanding Context and Time

RNNs and their variants (LSTMs, GRUs) are fundamental to systems that rely on understanding sequential relationships, enabling fluid, context-aware interactions.



Language Translation

Models like Google Translate use sequence-to-sequence RNNs to interpret full sentences and generate syntactically correct translations.



Speech Recognition

They convert spoken audio waveforms into text, powering virtual assistants like Siri and Alexa, by analyzing temporal acoustic patterns.

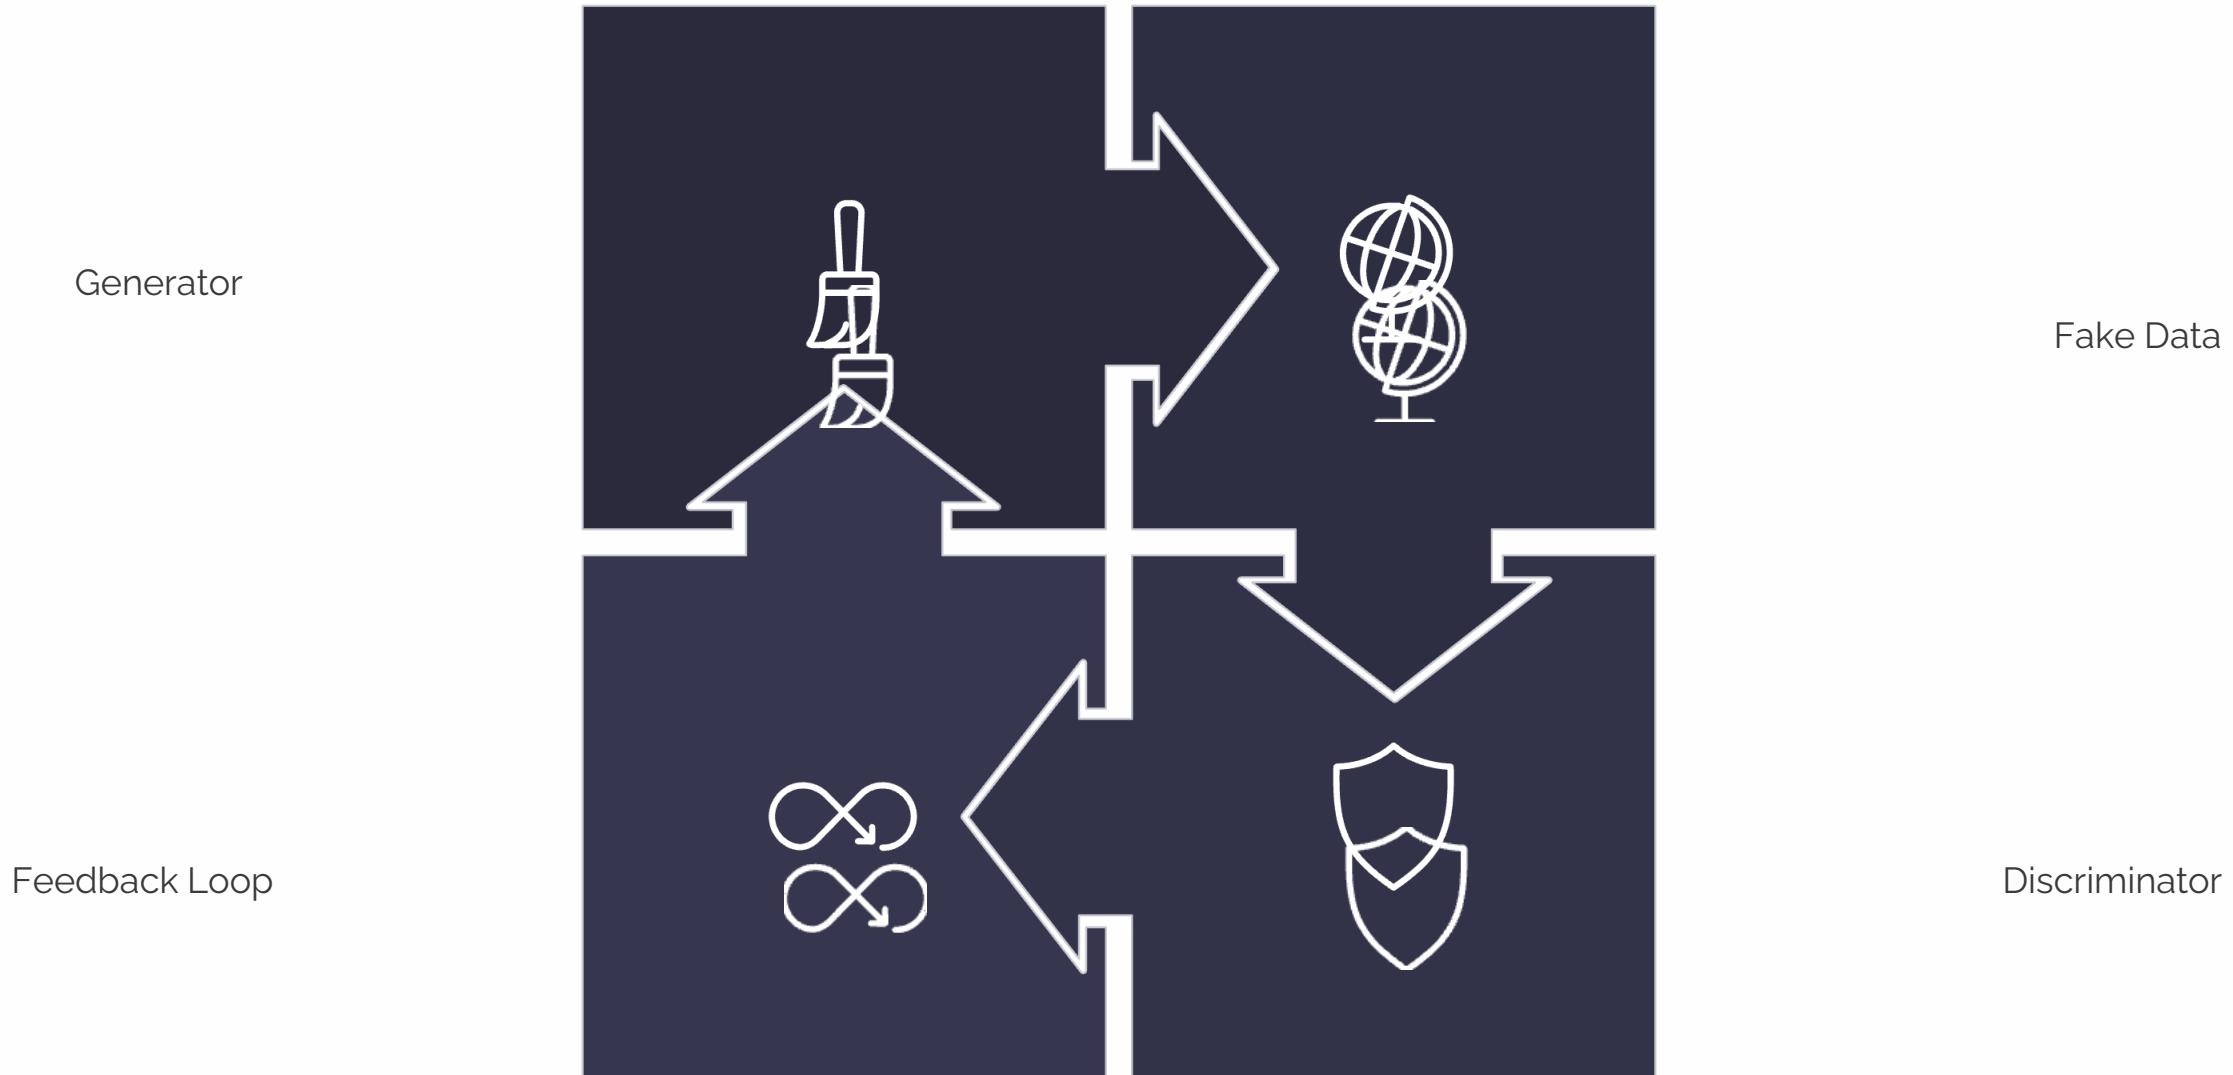


Time Series Prediction

RNNs analyze historical data—such as stock movements or weather patterns—to predict future values based on past trends.

Generative Adversarial Networks (GANs)

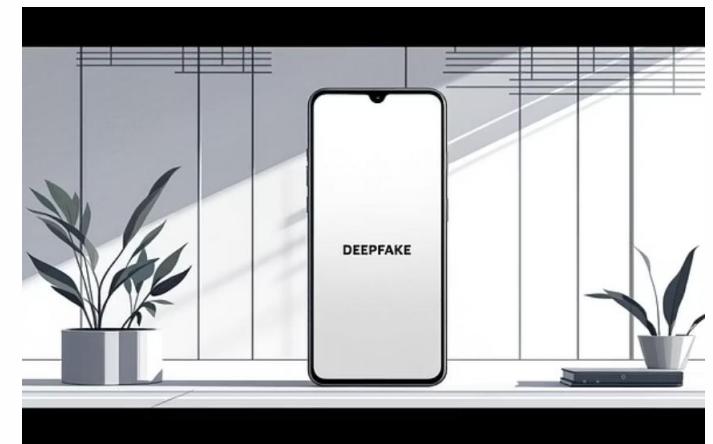
GANs represent a revolutionary approach to generative modeling, training two neural networks simultaneously in a competitive zero-sum game.



The Goal: The Generator aims to create perfect fakes, while the Discriminator aims to become a perfect detector. This adversarial process drives both models toward highly sophisticated performance.

GAN Applications: Creating the Unreal

From stunningly realistic synthetic images to highly effective data augmentation, GANs are changing how we interact with and create digital content.



→ Image Synthesis

Creating entirely novel and photorealistic data, such as faces (e.g., StyleGAN), environments, or objects.

→ Data Augmentation

Generating synthetic variations of existing training data to improve the robustness and generalization of other ML models.

→ Art and Creativity

Powering new tools for AI-driven art generation and artistic style transfer.

Conclusion & Future Directions

Specialization and Hybrid Systems

Aspect	CNN	RNN	GAN
Data Type	Images, Spatial Data	Sequences, Time Series	Any (Generation)
Core Function	Spatial Feature Extraction	Temporal Dependency & Context	Adversarial Synthesis
Example Output	Classification, Detection	Prediction, Text Generation	Synthetic Data, Art

The Future is Integrated:

Autonomous Driving: CNN for real-time vision processing + RNN for predicting future trajectory based on temporal data.

AI Image Captioning: CNN extracts features from the image + RNN generates a coherent descriptive sentence.